

# 1 Definitions

## 1.1 Notation

We will use  $\delta(v)$  to refer to the set of edges adjacent to  $v$ . We identify half-edges in a graph  $G = (V, E)$  with a pair  $(e, v)$  where  $e \in \delta(v)$ , referring to the half of  $e$  adjacent to  $v$ .

We use  $\{\{\star\}\}$  to refer to multisets, and the following notation for multisubsets of a set  $A$ :

- $A^{\{\{k\}\}}$  refers to the set of all multisubsets of  $A$  of size  $k \in \mathbb{N}$ ,
- $A^{\{\{\mathbb{N}\}\}}$  refers to the set of all finite multisubsets of  $A$ ,
- $A^{\{\{\star\}\}}$  refers to the set of all multisubsets of  $A$ .

## 1.2 Common and uncommon assumptions

**T-hop neighbourhoods:** we define a *T-hop centered neighbourhood* as a *centered graph of radius T* using the definitions from page 4 of [NS93]. Note that, in the case of subgraphs, the  $T$ -hop neighbourhood of  $v$  may be different from the subgraph induced by all nodes that have distance at most  $T$  from  $v$ .

**Local model:** we work in the *deterministic LOCAL* model with input labels. No assumptions are made about the number of input or output labels. We assume the nodes are aware of the exact number  $n$  of nodes in the input graph, but not of any information about the maximum degree  $\Delta$ . Because of this information, we equivalently describe a LOCAL algorithm as a  $T(n)$ -round communication algorithm with unbounded message size, or as a *possibly uncomputable* function from  $T(n)$ -hop centered neighbourhoods to output labels. Additionally, we assume there is a finite integer  $c$  such that every node is assigned a unique ID in the set  $\{1, \dots, n^c\}$ , but the nodes are not aware of the value of  $c$ .

**Solvability:** we say that a problem is *weakly unsolvable* if there are finitely many unsolvable instances for it. We will treat weakly unsolvable problems as solvable problems by adding one output label  $U$  and requiring the problem to output  $U$  on all nodes if the instance is unsolvable; this requires only constant time, so it does not affect asymptotic complexity. We call the maximum diameter of an unsolvable instance the *solvability horizon* of the weakly unsolvable problem. We say that a problem is *strongly unsolvable* if there are infinitely many unsolvable instances for it.

**LCL problems:** we define LCL problems as tuples  $\Pi = (\Sigma_{\text{in}}, \Sigma_{\text{out}}, r, \mathcal{C})$  where:

- $\Sigma_{\text{in}}$  is a set of input labels,
- $\Sigma_{\text{out}}$  is a set of output labels,
- $r$  is a finite integer, called the *radius* of  $\Pi$ , and
- $\mathcal{C}$  is a *finite* set of  $r$ -hop centered neighbourhoods, where each node is labelled by a pair in  $\Sigma_{\text{in}} \times \Sigma_{\text{out}}$ .

Though we haven't explicitly stated this in the definition, WLOG we can and will also assume that  $\Sigma_{\text{in}}$  is finite (else the problem would be strongly unsolvable) and that  $\Sigma_{\text{out}}$  is finite (any labels that don't appear in  $\mathcal{C}$  will not appear in any valid solution and can be ignored).

### 1.3 New-ish definitions

**Hanging trees:** *connected* tree graphs containing a finite number of incomplete half-edges, called *hooks*. Specifically, we will call hanging trees with one hook *ornaments* and hanging trees with two hooks *tinsels*.

**Class:** given an ornament  $T$  and a node-edge checkable problem  $\Pi$ , we define the *class* of  $T$  to be the set  $\mathfrak{Cl}(T) \subseteq \Sigma_{\text{out}}$  containing exactly the labellings of the hook that can be completed to a valid labelling for  $T$  according to  $\Pi$ . Up to the changed role of nodes and edges, this is the *class* definition from Section 3 of [CP17].

### 1.4 Actually new definitions

**Finitely Represented Configuration:** given an alphabet of symbols  $\Sigma$ , we call a *finitely represented configuration* of  $\Sigma$  a pair of a *finite* multisubset of  $\Sigma$  (the *requirement*) and a *finite* subset of  $\Sigma$  (the *filler*). We denote the set of finitely represented configurations in  $\Sigma$  as

$$\Sigma^{FSL} := \Sigma^{\{\{\mathbb{N}\}\}} \times [\Sigma]^{<\omega}$$

**Extended Finite State Labelling problems:** we define Extended FSL problems as tuples  $\Pi = (\Sigma_{\text{in}}, \Sigma_{\text{out}}, f, \mathcal{N}, \mathcal{E})$  where:

- $\Sigma_{\text{in}}$  is a *finite* set of input labels,
- $\Sigma_{\text{out}}$  is a *finite* set of output labels,
- $f : \mathbb{N} \rightarrow 2^{\Sigma_{\text{in}}}$  is a *computable* function describing which input labels are allowed to appear for a node of degree  $d \in \mathbb{N}$  which is surjective on some cover of  $\Sigma_{\text{in}}$  (in other words, each label in  $\Sigma_{\text{in}}$  appears in some set in  $f(\mathbb{N})$ ),
- $\mathcal{N} : \Sigma_{\text{in}} \rightarrow 2^{(\Sigma_{\text{out}})^{FSL}}$  is a function assigning to each input label a set of *finitely represented configurations* of  $\Sigma_{\text{out}}$ , and
- $\mathcal{E}$  is a set of pairs of elements of  $\Sigma_{\text{out}}$ .

An *instance* for  $\Pi$  is a graph  $G = (V, E)$  together with a labelling function  $l : V \rightarrow \Sigma_{\text{in}}$  such that  $l(v) \in f(\deg(v))$ , that is, the labelling is coherent with the “allowed” labels for each degree. A *solution* for this instance is a labelling  $s$  of the half-edges of  $G$  such that:

- $\forall v \in V$ , let  $S(v) := \{\{s(e, v) \mid e \in \delta(v)\}\}$ ; then there exists  $(R, F) \in \mathcal{N}(l(v))$  such that  $R \subseteq S(v)$  and  $(S(v) \setminus R) \subseteq F^{\{\star\}}$ , that is to say, every symbol in the requirement appears exactly once in  $S(v)$ , and every other symbol of  $S(v)$  is a filler symbol, and
- $\forall e = \{u, v\} \in E$  we have  $\{s(e, u), s(e, v)\} \in \mathcal{E}$ .

**Finite State Labelling problems:** we define FSL problems as tuples  $\Pi = (\Sigma_{\text{out}}, \mathcal{N}, \mathcal{E})$  where:

- $\Sigma_{\text{out}}$  is a *finite* set of output labels,
- $\mathcal{N} \subseteq 2^{(\Sigma_{\text{out}})^{FSL}}$  is a *finite* set of finitely represented configurations of  $\Sigma_{\text{out}}$ ,
- $\mathcal{E}$  is a set of pairs of elements of  $\Sigma_{\text{out}}$ .

An *instance* for  $\Pi$  is any graph; a solution for this instance is a labelling  $s$  of the half-edges of  $G$  such that:

- $\forall v \in V$ , let  $S(v) := \{\{s(e, v) \mid e \in \delta(v)\}\}$ ; then there exists  $(R, F) \in \mathcal{N}$  such that  $R \subseteq S(v)$  and  $(S(v) \setminus R) \subseteq F^{\{\star\}}$ , that is to say, every symbol in the requirement appears exactly once in  $S(v)$ , and every other symbol of  $S(v)$  is a filler symbol, and

- $\forall e = \{u, v\} \in E$  we have  $\{\{s(e, u), s(e, v)\}\} \in \mathcal{E}$ .

**Minimum degree of a label:** we define the *minimum degree* of an input label  $\chi \in \Sigma_{\text{in}}$  as  $\deg_{\min}(\chi) := \min \{d \in \mathbb{N} \mid \chi \in f(d)\}$  (the minimum degree for which that input label is allowed). This set is never empty. Generally, we can assume that  $\forall (R, F) \in \mathcal{N}(\chi)$  we have  $|R| = \deg_{\min}(\chi)$ , since a configuration with a longer requirement would be unattainable by a finite number of degrees (which can be encoded by putting them in separate input classes) and any configuration with a shorter requirement implicitly requires using a number of filler symbols, which can be encoded by adding all combinations of filler symbols to the requirements multiset.

**Width:** we define the *width* of a EFSL problem  $\Pi$  as the maximum length of one of its requirements: accounting for hidden requirements encoded in the degree, we get

$$\max \left\{ \max_{\chi \in \Sigma_{\text{in}}} \deg_{\min}(\chi), \max \{|R| \mid \exists \chi \in \Sigma_{\text{in}}, \exists F \subset \Sigma_{\text{out}} : (R, F) \in \mathcal{N}(\chi)\} \right\}$$

For FSL problems we define it as  $\max \{|R| \mid \exists F \subset \Sigma_{\text{out}} : R, F \in \mathcal{N}\}$ .

**Height:** we define the *height* of an EFSL problem as  $h(\Pi) := \max \{|\mathcal{N}(\chi)| \mid \chi \in \Sigma_{\text{in}}\}$  (the maximum number of configurations any node can be in). For FSL problems we define it as  $|\mathcal{N}|$ .

**Restricted (E)FSL:** we define the *restriction* of an (E)FSL the *node-edge checkable* problem with input labels obtained by replacing every finitely represented configuration  $(R, F)$  with the configurations  $\{R \cup \{\{x\}\} \mid x \in F\}$ .

## 2 List of results

1. Extended FSL = FSL (up to changing the number of labels) input/degree info in the structure and output labels
2. Node-edge checkable problems  $\subseteq$  FSL
3. Weak and strong solvability are decidable on FSL

## 3 List of possible results

1. If a FSL problem is known to be  $\Omega(\log n)$  on trees, its complexity on trees is decidable and an efficient algorithm can be found. (This holds because of Gustav's proof which should be stronger, it's in here because I still haven't written out all the details of the FSL based one)
2. (Belief I haven't been able to fully prove) For each FSL problem  $\Pi$  there is a value  $d$  that only depends on the problem description (specifically, number of labels, height and width) such that there is a worst-case family of graphs for  $\Pi$  of maximum degree  $d$ . My theory is either width+1 or width·height (also possible: width+height, using one of each filler). Latter is required for solvability.
3. (Investigating) A form of the Round Elimination procedure applies to FSL problems.

## 4 Proofs

### 4.1 Equivalence of EFSL and FSL

In this section, we show that the EFSL and FSL can encode the same problems.

**Theorem 4.1.** *For each FSL problem  $\Pi$  that has a LOCAL algorithm  $A$  with complexity  $T(n)$ , there is an EFSL problem  $\Pi'$  which has a LOCAL algorithm  $A'$  with complexity  $O(T(n))$ .*

*Proof.* Let  $\Pi = (\Sigma_{\text{out}}, \mathcal{N}, \mathcal{E})$ . We define  $\Pi' = (\Sigma'_{\text{in}}, \Sigma'_{\text{out}}, f', \mathcal{N}', \mathcal{E}')$  as follows:

- $\Sigma'_{\text{in}} = \{\star\}$ ,
- $\Sigma'_{\text{out}} = \Sigma_{\text{out}}$ ,
- $f'$  is the constant function  $f(n) = \{\star\}$ ,
- $\mathcal{N}'(\star) = \mathcal{N}$ ,
- $\mathcal{E}' = \mathcal{E}$ .

It is trivial to observe that this is the same problem, and can be solved with the same complexity as  $\Pi$ .  $\square$

**Theorem 4.2.** *For each EFSL problem  $\Pi$  that has a LOCAL algorithm  $A$  with complexity  $T(n)$ , there is an FSL problem  $\Pi'$  which has a LOCAL algorithm  $A'$  with complexity  $O(T(n))$ .*

This is the more complex of the two proofs, so we provide an overview before getting into the technical details: we encode the  $k$ -th input label as a line of degree 3 nodes of length  $k$ , and add leaves as needed to make sure every node of the original graph is of degree  $\geq 4$ . This at most multiplies the number of nodes by  $3|\Sigma_{\text{in}}|$ , and every node can in finite time compute its own input label by looking at its  $(|\Sigma_{\text{in}}| + 1)$ -hop neighbourhood. To make sure

*Proof.* Let  $\Pi = (\Sigma_{\text{in}}, \Sigma_{\text{out}}, f, \mathcal{N}, \mathcal{E})$ . Up to a bijection we can identify the set  $\Sigma_{\text{in}}$  with the set  $\{1, \dots, k\}$ . We define  $\Pi' = (\Sigma'_{\text{out}}, \mathcal{N}', \mathcal{E}')$  as:

- $\Sigma'_{\text{out}} = \Sigma_{\text{out}} \cup \Sigma_I$ , where  $\Sigma_I := \{(I, n) \mid n \in \Sigma_{\text{in}}\} \cup \{(I, L), (I, Err), Err\}$ ,
- $\mathcal{E}' = \mathcal{E} \cup \{\{\{X, X\}\} \mid X \in \Sigma_I\}$  (every edge containing the new input labels is symmetric).

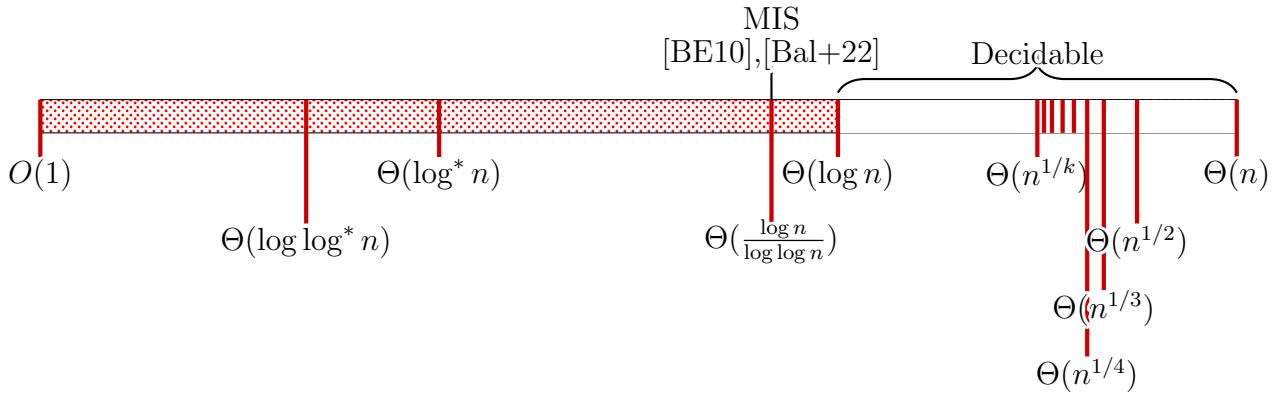
$\mathcal{N}'$  has a more complex formula, so we provide an algorithm to derive it from  $f$  and  $\mathcal{N}$  instead:

1.

$\square$

## 5 What do we know about the complexity landscape?

### 5.1 Trees



### 5.2 General graphs

## Bibliography

- [NS93] Moni Naor and Larry Stockmeyer. “What can be computed locally?” In: *Proceedings of the twenty-fifth annual ACM symposium on Theory of Computing*. STOC ’93. New York, NY, USA: Association for Computing Machinery, June 1, 1993, pp. 184–193. ISBN: 978-0-89791-591-5. DOI: 10.1145/167088.167149. URL: <https://www.wisdom.weizmann.ac.il/~naor/PAPERS/lcl.pdf> (visited on 08/11/2025).
- [BE10] Leonid Barenboim and Michael Elkin. “Sublogarithmic Distributed MIS Algorithm for Sparse Graphs Using Nash-Williams Decomposition”. In: *Distributed Computing* 22.5 (Aug. 1, 2010), pp. 363–379. ISSN: 1432-0452. DOI: 10.1007/s00446-009-0088-2. URL: <https://doi.org/10.1007/s00446-009-0088-2> (visited on 09/16/2025).
- [CP17] Yi-Jun Chang and Seth Pettie. “A Time Hierarchy Theorem for the LOCAL Model”. In: 2017 IEEE 58th Annual Symposium on Foundations of Computer Science (FOCS). Berkeley, CA, Oct. 2017, pp. 156–167. ISBN: 978-1-5386-3464-6. DOI: 10.1109/FOCS.2017.23. URL: <http://ieeexplore.ieee.org/document/8104055/> (visited on 08/11/2025).
- [Bal+22] Alkida Balliu et al. “Distributed  $\Delta$ -Coloring Plays Hide-and-Seek”. In: *Proceedings of the 54th Annual ACM SIGACT Symposium on Theory of Computing*. STOC 2022. New York, NY, USA: Association for Computing Machinery, June 10, 2022, pp. 464–477. ISBN: 978-1-4503-9264-8. DOI: 10.1145/3519935.3520027. URL: <https://doi.org/10.1145/3519935.3520027> (visited on 09/16/2025).